

A4 + A5

Type

Homework

Write-Ahead-Log (WAL)

A file where insert, update, delete information is written. it is fail safe.

1. Return the current LSN of your WAL file

```
SELECT pg_current_wal_lsn();
```

Data Output		Messages	Notifications
<div><div></div><div></div><div></div><div></div><div></div></div>		<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
	pg_current_wal_lsn pg_lsn		
1	0/1F6E780		

0/ is WAL segment, 1F6E780 is offset within the segment.

2. Run the transaction on next slide. (Make sure that the Pks do NOT yet exist in your table lesson.)

```
Begin;
update student set s_balance = s_balance - 3 where s_username= 'Rose';
update teacher set t_payment = t_payment + 3 where t_id = 1;
select t_payment from teacher where t_id = 1;
INSERT INTO lesson(t_id, lesson_time, s_username, subjectcode)
VALUES (1, '2024-03-03 05:22:12.000000', 'Rose', 'EN');
INSERT INTO lesson(t_id, lesson_time, s_username, subjectcode)
VALUES (1, '2024-03-04 05:22:12.000000', 'Rose', 'EN');
INSERT INTO lesson(t_id, lesson_time, s_username, subjectcode)
VALUES (1, '2024-03-05 05:22:12.000000', 'Rose', 'EN');
Commit;
```

3. Return the current LSN of your WAL file

Data Output		Messages	Notifications
<div><div></div><div></div><div></div><div></div><div></div></div>		<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
	pg_current_wal_lsn pg_lsn		
1	0/1F711A8		

4. Dump the WAL from the first to the second LSN

```
pg_waldump -s 0/1F6E780 -e 0/1F711A8 -p "D:/postgreSQL/17/data/pg_wal"
```

```
PS C:\Users\barba> pg_waldebug -s 0/1F6E780 -e 0/1F711A8 -p "D:/postgresql/17/data/pg_wal"
rmgr: Heap len (rec/tot): 65/ 853, tx: 805, lsn: 0/01F6E780, prev 0/01F6E748, desc: HOT_UPDATE old_xmax: 805, old_off: 7, old_infobits:
[ ] flags: 0x00, new_xmax: 0, new_off: 9, blkref #0: rel 1663/16478/16493 blk 0 FPM
rmgr: Heap len (rec/tot): 65/ 6461, tx: 805, lsn: 0/01F6EAD8, prev 0/01F6E780, desc: HOT_UPDATE old_xmax: 805, old_off: 20, old_infobits:
[ ] flags: 0x00, new_xmax: 0, new_off: 21, blkref #0: rel 1663/16478/16499 blk 0 FPM
rmgr: Heap len (rec/tot): 54/ 830, tx: 805, lsn: 0/01F70430, prev 0/01F6EAD8, desc: INSERT off: 12, flags: 0x00, blkref #0: rel 1663/164
78/16482 blk 0 FPM
rmgr: Standby len (rec/tot): 54/ 54, tx: 0, lsn: 0/01F70770, prev 0/01F70430, desc: RUNNING_XACTS nextXid 806 latestCompletedXid 804 old
estRunningXid 805; 1 xacts: 805
rmgr: Btree len (rec/tot): 53/ 429, tx: 805, lsn: 0/01F707A8, prev 0/01F70770, desc: INSERT_LEAF off: 4, blkref #0: rel 1663/16478/16517
blk 1 FPM
rmgr: Btree len (rec/tot): 53/ 429, tx: 805, lsn: 0/01F70958, prev 0/01F707A8, desc: INSERT_LEAF off: 8, blkref #0: rel 1663/16478/16533
blk 1 FPM
rmgr: Heap len (rec/tot): 54/ 54, tx: 805, lsn: 0/01F70B08, prev 0/01F70958, desc: LOCK xmax: 805, off: 21, infobits: [LOCK_ONLY, KEYSH
R_LOCK], flags: 0x00, blkref #0: rel 1663/16478/16499 blk 0
rmgr: Heap len (rec/tot): 59/ 1801, tx: 805, lsn: 0/01F70B40, prev 0/01F70B08, desc: LOCK xmax: 805, off: 1, infobits: [LOCK_ONLY, KEYSHR
_LOCK], flags: 0x00, blkref #0: rel 1663/16478/16505 blk 0 FPM
rmgr: Heap len (rec/tot): 54/ 54, tx: 805, lsn: 0/01F70F88, prev 0/01F70B40, desc: LOCK xmax: 805, off: 9, infobits: [LOCK_ONLY, KEYSHR
_LOCK], flags: 0x00, blkref #0: rel 1663/16478/16493 blk 0
rmgr: Heap len (rec/tot): 79/ 79, tx: 805, lsn: 0/01F70FC0, prev 0/01F70F88, desc: INSERT off: 13, flags: 0x00, blkref #0: rel 1663/164
78/16482 blk 0
rmgr: Btree len (rec/tot): 72/ 72, tx: 805, lsn: 0/01F71010, prev 0/01F70FC0, desc: INSERT_LEAF off: 5, blkref #0: rel 1663/16478/16517
blk 1
rmgr: Btree len (rec/tot): 72/ 72, tx: 805, lsn: 0/01F71058, prev 0/01F71010, desc: INSERT_LEAF off: 9, blkref #0: rel 1663/16478/16533
blk 1
rmgr: Heap len (rec/tot): 79/ 79, tx: 805, lsn: 0/01F710A0, prev 0/01F71058, desc: INSERT off: 14, flags: 0x00, blkref #0: rel 1663/164
78/16482 blk 0
rmgr: Btree len (rec/tot): 72/ 72, tx: 805, lsn: 0/01F710F0, prev 0/01F710A0, desc: INSERT_LEAF off: 6, blkref #0: rel 1663/16478/16517
blk 1
rmgr: Btree len (rec/tot): 72/ 72, tx: 805, lsn: 0/01F71138, prev 0/01F710F0, desc: INSERT_LEAF off: 10, blkref #0: rel 1663/16478/16533
blk 1
rmgr: Transaction len (rec/tot): 34/ 34, tx: 805, lsn: 0/01F71180, prev 0/01F71138, desc: COMMIT 2025-03-26 03:10:29.937503 Georgian Standard
Time
rmgr: Transaction len (rec/tot): 34/ 34, tx: 805, lsn: 0/01F71180, prev 0/01F71138, desc: COMMIT 2025-03-26 03:10:29.937503 Georgian Standard
Time
```

1. Hot Updates

- `old_xmax: 805` - This is **transaction ID** that invalidated the **old row**.
- `old_off: 20` - This represents **offset** of the old row within the page.
- `new_xmax: 0` - This is **transaction ID** that (will or not) delete the new row.
- `new_off: 21` - This is the **offset** of the newly inserted row within the same page.

2. why inserts into 4 different btrees?

- `INSERT` - Insert new row in **Heap**.
- `INSERT_LEAF` - Insert new index in **BTree**.

BTree data structures are added to **primary key** and **unique constraints**.

In our case tuple, `t_id` and `lesson_time`, form primary composite key (1 **BTree**).

We also have unique constraint on tuple, `t_id` and `s_username`, forming second **BTree**.

3. why inserts into the btree leaves?

Since the **actual index data is stored in leaf nodes**, PostgreSQL only logs `INSERT_LEAF`.

Dirty Read Anomaly in PostgreSQL?

reads data written by a concurrent uncommitted transaction.

Transaction 1

```
Begin;
drop table if exists transaction_log;
CREATE TEMPORARY TABLE transaction_log (message_text varchar(50), t_payment_value INT);
INSERT INTO transaction_log (message_text, t_payment_value) select 'Payment_amount',t_payment FROM teacher WHERE t_id = 1;
UPDATE teacher SET t_payment = 0 WHERE t_id = 1;

SELECT pg_sleep(30);

INSERT INTO transaction_log (message_text, t_payment_value) select 'after_reset_balance',t_payment FROM teacher
WHERE t_id = 1;
rollback;
```

Transaction 2

```
Begin;
drop table if exists transaction_log;
CREATE TEMPORARY TABLE transaction_log (t_payment_value INT);
```

```
INSERT INTO transaction_log (t_payment_value) select t_payment FROM teacher WHERE t_id = 1;
COMMIT;
```

Consistent result would be:

`t_payment=0` shouldn't have been for `t_id=1` it should have had its original value.

Inconsistent result ("dirty read") would be:

it reads uncommitted data which is `t_payment_value = 0`, thus causing dirty read.

```
agency=# SELECT * FROM transaction_log;
 t_payment_value
-----
                0
(1 row)
```

Does it make any difference for the read in session 2 if the transaction in session 1 commits or rollbacks?

if it would've rolledback it would've had its original value and for the commit it would have the value of 0.

Non-Repeatable Read Anomaly

reads the same row twice but gets different data each time. in repeatable-read and serializable in postgresSQL.

Transaction 1

```
BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
SELECT s_username FROM student WHERE s_balance < 5;
SELECT pg_sleep(20);
SELECT s_username FROM student WHERE s_balance < 5;
COMMIT;
```

```
agency=# SELECT s_username FROM student WHERE s_balance < 5;
 s_username
-----
Witzi
Wolf
Darthvader
(3 rows)

agency=# COMMIT;
COMMIT
agency=# BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
BEGIN
agency=# SELECT s_username FROM student WHERE s_balance < 5;
 s_username
-----
Witzi
Wolf
Darthvader
(3 rows)
```

Transaction 2

```
BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
UPDATE student SET s_balance = s_balance + 5 WHERE s_username IN
(SELECT s_username FROM student WHERE s_balance < 5 LIMIT 1);
COMMIT;
```

```
BEGIN
agency==# UPDATE student SET s_balance = s_balance + 5 WHERE s_username IN
agency--# (SELECT s_username FROM student WHERE s_balance < 5 LIMIT 1);
UPDATE 1
agency==# COMMIT;
COMMIT
agency=# select s_username from student where s_balance < 5;
 s_username
-----
Darthvader
(1 row)
```

Phantom Read Anomaly in PostgreSQL?

re-executes a query and sees new rows inserted by another transaction, phantom read in postgresQL in repeatable read and serializable. In other only serializable.

Transaction 1

```
BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
SELECT count(*) FROM teacher WHERE t_payment <= 5;
SELECT pg_sleep(20);
SELECT count(*) FROM teacher WHERE t_payment <= 5;
COMMIT;
```

```
agency=# BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
BEGIN
agency==# SELECT count(*) FROM teacher WHERE t_payment <= 5;
 count
-----
      18
(1 row)

agency==# SELECT pg_sleep(20);
 pg_sleep
-----
(1 row)

agency==# SELECT count(*) FROM teacher WHERE t_payment <= 5;
 count
-----
      18
(1 row)

agency==# COMMIT;
COMMIT
```

Transaction 2

```
BEGIN;
INSERT INTO public.teacher (t_id, t_name, t_mail, t_postalcode, t_dob, t_gender, t_education, t_remark, t_payment)
VALUES (25, 'lamara', 'lamar@gmail.com', 8000, '1880-10-10', 'f', 'Bachelor', NULL, 0);
COMMIT;
```

```
agency=# SELECT count(*) FROM teacher WHERE t_payment <= 5;
 count
-----
      19
(1 row)
```

2PL

Given is the following schedule. How would it work under 2PL?

```
T1:Read(A) T2:Read(A) T1:Write(A) T2:Write(A) T1:Read(B) T2:Read(B) T1:Write(B)
T2:Write(B)
```

What is the result?

Operation	T1	T2
1	BEGIN	
2	S-lock on A	BEGIN
3		S-lock on A
4	can't put E-lock on A in T1, because of having to wait for T2	
5		can't put E-lock on A in T2, because of having to wait for T1
6	Deadlock	Deadlock

Deadlock

1. Write the code for T1, T2 and T3 and let them run parallel in 3 sessions.

Transaction 1

```
BEGIN;
UPDATE student SET s_balance = s_balance - 3 WHERE s_username = 'Mickey';
SELECT pg_sleep(40);
UPDATE teacher SET t_payment = t_payment + 3 WHERE t_id = 1;
COMMIT;
```

Transaction 2

```
BEGIN;
UPDATE teacher SET t_payment = t_payment + 3 WHERE t_id = 1;
SELECT pg_sleep(35);
UPDATE student SET s_balance = s_balance - 3 WHERE s_username = 'Rose';
COMMIT;
```

Transaction 3

```
BEGIN;
UPDATE student SET s_balance = s_balance - 3 WHERE s_username = 'Rose';
SELECT pg_sleep(30);
UPDATE student SET s_balance = s_balance - 3 WHERE s_username = 'Mickey';
UPDATE teacher SET t_payment = t_payment + 6 WHERE t_id = 2;
COMMIT;
```

2. Which transaction is rolled back, which commit in what sequence?

First Transaction 2 rollbacked

```
ERROR:  deadlock detected
DETAIL:  Process 11772 waits for ShareLock on transaction 822; blocked by process 10072.
Process 10072 waits for ShareLock on transaction 820; blocked by process 7184.
Process 7184 waits for ShareLock on transaction 821; blocked by process 11772.
HINT:  See server log for query details.
CONTEXT:  while updating tuple (0,9) in relation "student"
agency=!# COMMIT;
ROLLBACK
```

Second Transaction 1 committed

```
agency=# SELECT pg_sleep(40);
pg_sleep
-----

(1 row)

agency=# UPDATE teacher SET t_payment = t_payment + 3 WHERE t_id = 1;
UPDATE 1
agency=# COMMIT;
COMMIT
```

Third Transaction 3 committed

```
agency=# UPDATE student SET s_balance = s_balance - 3 WHERE s_username = 'Rose';
UPDATE 1
agency=# SELECT pg_sleep(30);
pg_sleep
-----

(1 row)

agency=# UPDATE student SET s_balance = s_balance - 3 WHERE s_username = 'Mickey';
UPDATE 1
agency=# UPDATE teacher SET t_payment = t_payment + 6 WHERE t_id = 2;
UPDATE 1
agency=# COMMIT;
COMMIT
```

3. What is the error message?

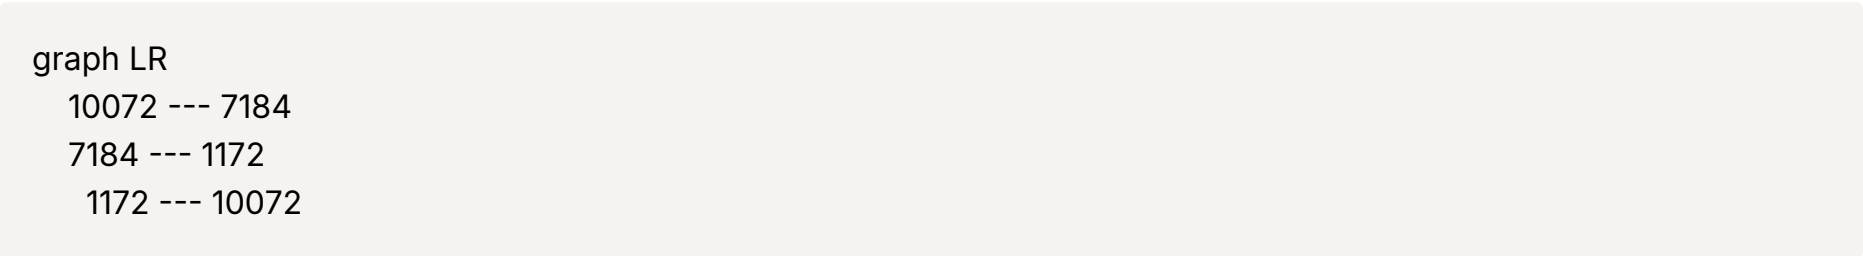
```
ERROR: deadlock detected
DETAIL: Process 11772 waits for ShareLock on transaction 822; blocked by process 10072.
Process 10072 waits for ShareLock on transaction 820; blocked by process 7184.
Process 7184 waits for ShareLock on transaction 821; blocked by process 11772.
HINT: See server log for query details.
CONTEXT: while updating tuple (0,9) in relation "student"
agency=!# COMMIT;
ROLLBACK
```

deadlock has been detected and second transaction has been aborted.

4. Explain the problem and the result.

the problem is that three transactions that run cause a deadlock, thus making one of the transactions abort and others to commit.

Waits-For-Graph



second transaction has caused the error it is the victim, it is the transaction that triggers the circuit, this happens because of pg_sleep, if it was ran separately as blocks T3 would’ve been aborted.

MVCC

```
T1 read(a) T1write(a) T2read(a) T2write(b) T1 read(b) T1write(c)
```

Assume that all writes are updates. Objects a, b and c already exist.
How does the execution under MVCC look like? Do the two transactions run through?

Given is the following schedule s:
T1 read(a) T1write(a) T2read(a) T2write(b) T1 read(b) T1write(c)
Assume that all writes are updates. Objects a,b,a nd c already exist.
How does the execution under MVCC look like? Do the two transactiobn run through?

T10	T11		Version	Value	Created_by	Deleted_by
Begin			A0	100	–	10
Read(a)			A1	150	10	0
Write(a)	Begin		B0	300	–	11
	Read(a)		B1	350	11	0
	Write(b)		C0	500	10	10
Read(b)	commit		C1	550	10	0
Write(c)						
commit						

Given is the following schedule s:
T1 read(a) T1 write(a) T2 read(a) T2 write(b) T1 read(b) T1 write (c)
Assume that object a and b exist. T2 write(b) is a delete and T1 write (c) is an insert. How does the execution under MVCC look like?

T1	T2		Version	Value	Created_by	Deleted_by
Begin			A0	100	–	1
read(a)			A1	150	1	0
write(a)	begin		B0	300	–	2
	read(a)		B1	350	2	0
	write(b)		C0	500	1	0
read(b)						
write(c)						

Write Skew anomaly

Creating table

```
CREATE TABLE personnel (
  name VARCHAR(200) PRIMARY KEY,
  on_duty BOOLEAN
);
```

```
INSERT INTO personnel (name, on_duty) VALUES
('Vanhelsing', TRUE),
('Bobby', TRUE),
('Charlieputh', TRUE),
('Davidthethird', FALSE),
('Eve', FALSE),
('Alice', TRUE),
('Bob', TRUE);
```

Transaction 1

```

BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
DO $$
DECLARE n_on_duty INTEGER;
BEGIN

    SELECT count(*) INTO n_on_duty FROM personnel WHERE on_duty = TRUE;

    IF n_on_duty > 2 THEN
        UPDATE personnel SET on_duty = FALSE WHERE name = 'Alice';
    END IF;
END $$;
SELECT pg_sleep(20);
COMMIT;

```

```

agency=# BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
BEGIN
agency== DO $$
agency$$$ DECLARE n_on_duty INTEGER;
agency$$$ BEGIN
agency$$$
agency$$$     SELECT count(*) INTO n_on_duty FROM personnel WHERE on_duty = TRUE;
agency$$$
agency$$$     IF n_on_duty > 2 THEN
agency$$$         UPDATE personnel SET on_duty = FALSE WHERE name = 'Alice';
agency$$$     END IF;
agency$$$ END $$;
DO
agency== SELECT pg_sleep(20);
pg_sleep
-----
(1 row)

agency== COMMIT;|

```

Transaction 2

```

BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
DO $$
DECLARE n_on_duty INTEGER;
BEGIN
    SELECT count(*) INTO n_on_duty FROM personnel WHERE on_duty = TRUE;

    IF n_on_duty > 2 THEN
        UPDATE personnel SET on_duty = FALSE WHERE name = 'Bob';
    END IF;
END $$;
SELECT pg_sleep(20);
COMMIT;

```

```

agency=# BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
BEGIN
agency== DO $$
agency$$$ DECLARE n_on_duty INTEGER;
agency$$$ BEGIN
agency$$$     SELECT count(*) INTO n_on_duty FROM personnel WHERE on_duty = TRUE;
agency$$$
agency$$$     IF n_on_duty > 2 THEN
agency$$$         UPDATE personnel SET on_duty = FALSE WHERE name = 'Bob';
agency$$$     END IF;
agency$$$ END $$;
DO
agency== SELECT pg_sleep(20);
pg_sleep
-----
(1 row)

agency== COMMIT;
COMMIT

```



```
agency=# SELECT count(*) INTO n_on_duty FROM personnel WHERE on_duty = TRUE;
SELECT 1
```

Running this queries in repeatable read won't change anything since `n_on_duty` is already set and this isolation level treats both `update` actions valid. we can only try `SERIALIZABLE`.

Delete from personnel

```
DELETE FROM personnel;
INSERT INTO personnel (name, on_duty) VALUES
('Vanhelsing', TRUE),
('Bobby', TRUE),
('Charlieputh', TRUE),
('Davidthethird', FALSE),
('Eve', FALSE),
('Alice', TRUE),
('Bob', TRUE);
```

replace `repeatable read` with `serializable`.

Second transaction throws an error.

```
agency=# COMMIT;
ERROR:  could not serialize access due to read/write dependencies among transactions
DETAIL:  Reason code: Canceled on identification as a pivot, during commit attempt.
HINT:  The transaction might succeed if retried.
agency=#
```

second transaction is not committed and `on_duty` doctors remain at least 4.

```
agency=# SELECT count(*) FROM personnel WHERE on_duty = 'True';
count
-----
      4
(1 row)
```