

# Lecture 6

▼ Type	Lecture
--------	---------

## NoSQL databases

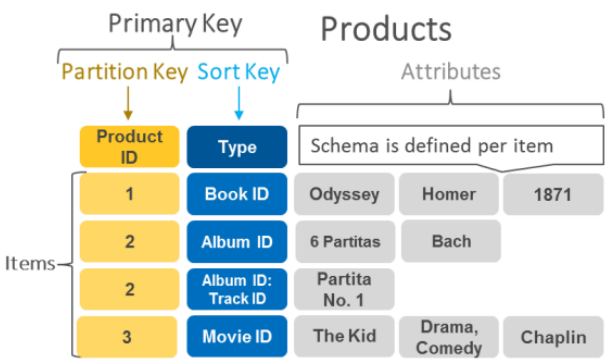
NoSQL includes key-value, document, wide-column, and graph databases.

- Designed for scalability, high performance, and flexible data models.

## Amazon DynamoDB

DynamoDB is not a classical key-value database, rather a key-value database with some wide-column database characteristics.

A key-value database is a non-relational database that uses a simple key-value method to store data. A key-value database stores data as a collection of key-value pairs in which a key serves as a unique identifier. Both keys and values can be arbitrary elements - from simple objects to complex composite objects.



Amazon Dynamo DB is a non-relational database. ... In DynamoDB, an element consists of a key and a flexible number of attributes (=values). There is no explicit restriction on the number of attributes associated with a single element, but the total size of an element, including all attribute names and values, cannot exceed 400 KB.

## Relational Database – Key-Value Database

Relational databases often use artificial primary keys and support schema-based searches.

In a relational database, when you create a table, you typically start by defining a primary key (PK)— an attribute that uniquely identifies each row in the table.

```
CREATE TABLE IF NOT EXISTS users (  
  user_id INT PRIMARY KEY AUTO_INCREMENT,  
  username VARCHAR(50),  
  email VARCHAR(100)  
);
```

## Key-Value Database

### Key in a Key-Value Database

Keys must be meaningful and descriptive because they are the only access method. Without knowing the key, the data is unreachable (no **WHERE** clauses or joins).

### Key-Value Databases

Key-Value (KV) databases are the simplest form of NoSQL systems. Data is stored as pairs: a unique key and its associated value.

### Important:

- No fixed structure (schema-free).
- Cannot search by value — only by key.
- No joins, constraints, or relationships between pairs.

- The application handles all logic related to how values are interpreted.
- Examples include Redis, Riak, and DynamoDB.

**Good Practices:**

- Make keys descriptive and unique.
- Use consistent naming conventions (e.g., `user:123:name` ).
- Avoid storing sensitive data in the key.

**Access to Data**

- Set / Put / Post      insert / update  
If the key does not exist yet, a pair will be inserted. If the key exists, the value will be overwritten.
- Get      reads the value of a pair
- Delete      deletes an entire key-value pair

put(key, value)

get(key, value) value=get(key);

delete(key)

---

**Use Case — Shopping Cart**

High write operations (add/update/remove).

**Important:** KV stores are ideal for temporary, frequently changing data.

**Use Case — Real-Time Counters**

Tracks scores or votes instantly.

**Important:** High speed and minimal conflict resolution required.

**Use Case: Storing Authentication or Session Information (Preferences, Personalized Marketing)**

Session ID or UUID is the key.

**Important:** Enables personalization by quickly fetching session data.

**Use Case: Sensor Data**

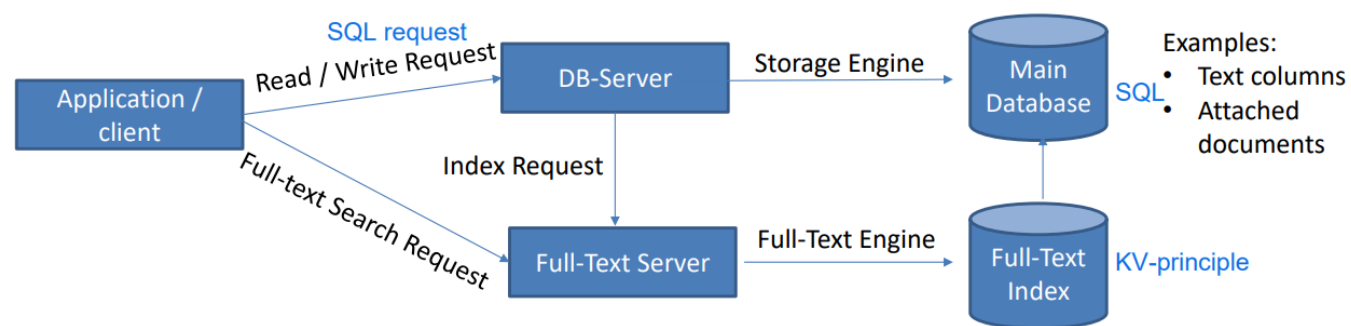
Periodic updates like temperature readings.

**Important:** Suited for IoT and monitoring systems.

**Use Case: Full-Text Index Concepts**

Full-text engines use key-value logic to map terms (keys) to documents (values).

**Important:** Often integrated with or run alongside main DBs.



When we say that a full-text index works like a key-value store, then what are the keys and what are the values?

keys are search terms, values are locations

1  
knirsch@htw-berlin.d

In general: Business cases with high velocity where the keys are enough to work with the data.

## Inverted Index

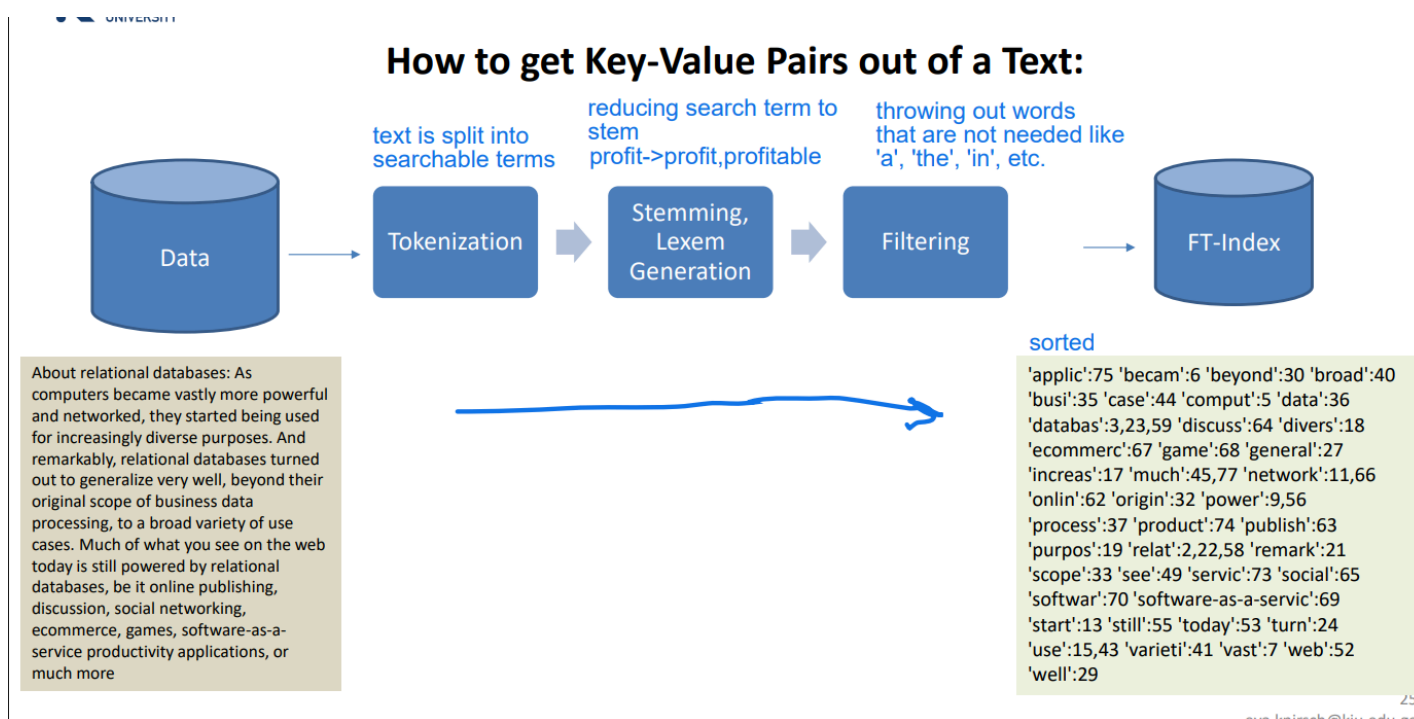
Each unique key (searchable term) is associated with a list of locations (pointers) where that term appears in the text or texts.

`tsvector` and GIN indexes build the full-text search.

**Important:** Efficient for indexing and searching large text data.

## Tokenization and Text Processing

Text is broken down into searchable parts (tokens, stems, lexemes). Filtering removes common words, improving search quality.



To get from text to ts-vector format:

**Tokenizer:** Parsing text into tokens

**Stemmer:** Converting tokens into stems (lexemes)

**Filter:** Eliminating stop words

26

## PostgreSQL Fulltext

Use `to_tsvector()` and `to_tsquery()` to search text columns. GIN indexes provide fast lookup for matching terms.

## PostgreSQL Fulltext

- PostgreSQL also includes fulltext indexer and fulltext search.
- PostgreSQL by default stores the fulltext index per row in the tables as extra key-value column of type `ts_vector`
- A GIN index (B-tree) on the `ts_vector` column then builds the fulltext index out of all `ts_vector` rows as fulltext index of the whole table or for multiple tables.

To perform a fulltext search:

Functions `to_tsvector()` and `to_tsquery` are used together with matching operator `@@`

Ad-hoc full text search:

```
SELECT * FROM table WHERE to_tsvector(text column) @@ to_tsquery('searchterm');
```

Full text search on existing column of type `ts_vector`

```
SELECT * FROM table WHERE ts_vector_column @@ to_tsquery('searchterm');
```

29