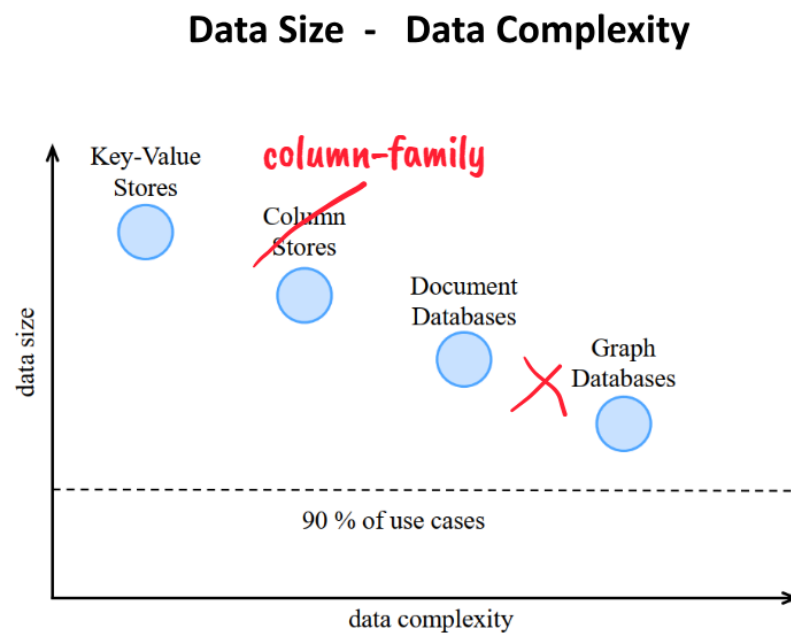


Lecture 14

Type Lecture

Data Size - Data Complexity



relational databases on the right of the x-axis, on X

Column-Family Databases

What is in a name?

Column Database / Column Store

Variation of a relational db

Column-Oriented Database / Columnar Database

Column-Family Database / Wide-Column Database

NoSQL database
P2P distributed
db

Row Format

- RDBMS are logically - two-dimensional: columns (attributes) and rows (tuples).
- However, the storage of the data happens one-dimensionally – in a file (heap file or B-Tree).
→ how to map from two dimensions to one dimension for storage?

teacherID	name	postal	email	DoB	gender	education	remark
41	Krawinkel	10045	krawinkel@wir_wissen_alles.xx	1978-09-15	f	Master	
42	Kaiser	10045	kaiser@zirkus.xx	2001-11-30	f	Ph.D	
43	Kern	10045	kern@immer_da.xx	1977-06-11	m	Ph.D	
44	Schuster	10004	schusterl@flieg.xx	1975-03-27	f	Master	
45	Newman	10010	new@kriech.xx	1995-11-20	m	Ph.D	

heap file storage

Header, 41, Krawinkel, 10045, krawinkel@wir_wissen_alles.xx, 1978-09-15, f, Master, Header, 42, Kaiser, 10045, kaiser@zirkus.xx, 2001-11-30, f, Ph.D, Header, 43, Kern, 10045, kern@immer_da.xx, 1977-06-11, m, Ph.D., Header, 44, Schuster 10004, schusterl@flieg.xx, 1975-03-27, f, Master, Header, 45, Newman, 10010, new@kriech.xx, 1995-11-20, m, Ph.D. **Header, 41, Krawinkel, 10047,..**

6
knirsch@htw-berlin.de

Column Store

teacherID	name	zip	rating
41	Krawinkel	10045	4
42	Kaiser	10045	5
43	Kern	10045	3
44	Schuster	10004	5
45	Neumann	10010	5

[41,42,43,44,45]
[Krawinkel,Kaiser,Kern,Schuster,Neumann]
[10045,10045,10045,10004,10010]
[4,5,3,5,5]

select * from teacher where teacherID = 44

How is this query executed?

array position

Why not store a TID with each column value?

Column Compression

profID	name	zip	rating
41	Krawinkel	10045	4
42	Kaiser	10045	5
43	Kern	10045	3
44	Schuster	10004	5
45	Neumann	10010	5

41,42,43,44,45;
Krawinkel,Kaiser,Kern,Schuster,Neumann;
10045,10045,10045,10004,10010;
4,5,3,5,5

- Columns often hold redundant values and / or a very limited number of different values. Those columns are suitable for compression.
- Popular compression method for columnar storage are bitmaps: each possible value gets a bitmap assigned. Each bitmap has same length as the column vector.
- Run-length encoding is used to further compact the bitmaps.

Bitmaps + Run-Length Encoding

- bitmap index on column rating with values {1,2,3,4,5}.
- The column vector has length 5 (5 tuples, Ids 41 - 45).
- Each rating value {1,2,3,4,5} gets a bitmap with length 5
- For each row, a 1 is filled into the bitmap if the row has that value; a 0 if it does not have the value.

- For each **Bitmap**

	1	2	3	4	5		rating
41	0	0	0	1	0	:	4
42	0	0	0	0	1		5
43	0	0	1	0	0		3
44	0	0	0	0	1		5
45	0	0	0	0	1		5

What advantage does compression have?

Run_length Encoding

	5	5	2,1	0,1	1,1,1,2
	5 zeros	5 zeros	2 zeros, 1 one, rest zeros	0 zeros, 1 one rest 2	1 zero, 1 one, 1 zero, 2 ones

11

eva.knirsch@kiu.edu.ge

for compaction. advantage of the compression is performance.

NoSQL

Column-Family Databases

1

13

RDBMS versus Column-Family Concepts

RDBMS:

schema
referential integrity
normalization --> minimize
redundancy --> consistency

join

ACID

"miniworld"

Column-Family:

P2P - token ring

schema
no referential integrity (no fks)

denormalized, redundant data
no joins --> the query needs to be
answered by one table

Designing for optimal storage
Sorting is a design decision

CAP - AP

https://cassandra.apache.org/doc/latest/cassandra/developing/data-modeling/data-modeling_rdbms.html

Defining Application Queries

keyspace lesson


Q1: Return all information about a given, specific teacher

John

Q2: Return the students of a given, specific teacher

Ana

Q3: Return a teacher list by given, specific postcode



How are the queries mapped on the components of the Cassandra data model:
→ a table for each query

- Rowkeys (=partition keys)
- column-families (attribute group)
- Columns (attribute – value)

controls the distribution onto the ring

Create Table Command for RDBMS (Postgres) and Column-Family (Cassandra)

Postgres:

```
CREATE TABLE IF NOT EXISTS public.teacher ((
  t_id integer NOT NULL DEFAULT
  nextval('teacher_tID_seq'::regclass),
  t_name character varying(30) COLLATE pg_catalog."default" NOT
  NULL,
  t_mail character varying(50) COLLATE pg_catalog."default" NOT
  NULL,
  t_postalcode integer NOT NULL,
  t_gender gender,
  t_education education NOT NULL,
  t_payment integer NOT NULL DEFAULT 0,
  CONSTRAINT teacher_pkey PRIMARY KEY (t_id)
))
```

Cassandra:

```
CREATE TABLE lesson.teachers (
  t_email text,
  t_name text,
  t_phone text,
  address frozen<address>,
  t_subjects set<text> ,
  t_education text,
  PRIMARY KEY (t_email)
  WITH comment = 'Q1. Find information about a teacher';
```

Table answers 1st query

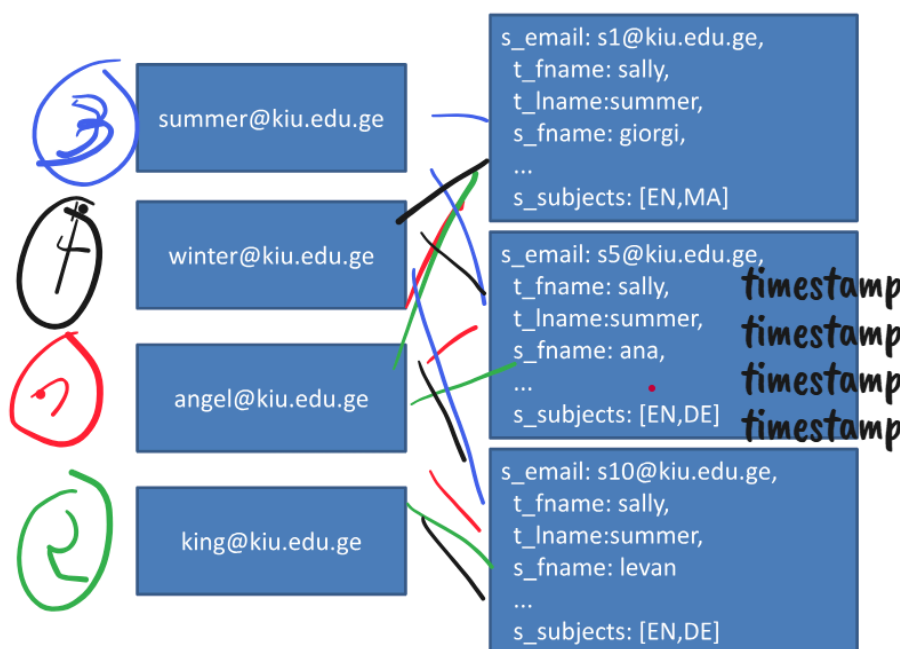
Table definition in Postgres and Cassandra looks similar. That is the problem.

Concepts and behavior are very different → data design needs to be approached differently, application development is different.

15

Partition key

Students



- Assume that these partition keys all are hashed to the same node
- Assume that the teachers have the same students
- What would the sorting in the memtable be?

26

knirsch@htw-berlin.de

flushed to disk

Use Case Examples

Storing user history of a service (e.g. streaming service, driving service)

```
CREATE TABLE history_by_user (  
  user_id UUID,  
  usage_timestamp TIMESTAMP, show_id TEXT,  
  title TEXT,  
  genre TEXT,  
  duration_watched INT,  
  location TEXT,  
  PRIMARY KEY ((user_id), view_timestamp usage_timestamp)  
) WITH CLUSTERING ORDER BY (usage_timestamp DESC);
```

Other use cases:

- Global Press Agencies storing their articles / press releases by date / topic
- Social media: feeds by followed persons per user